

Libraries

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression, Perceptron
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.linear_model import Perceptron
import seaborn as sns
import matplotlib.pyplot as plt
```

Load and Split Data set

```
dataset = pd.read_csv('Q2 Sentiment Analysis Dataset.csv',
encoding='latin-1', usecols=[0, 1, 2, 3])
class_names = [1, 3, 5, 'not_relevant']
X = dataset['text']
y = dataset['sentiment']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
dataset.head()
```

	id	sentiment	date \
0	623495523	1	Mon Dec 01 20:46:01 +0000 2014
1	623495527	1	Mon Dec 01 21:09:50 +0000 2014
2	623495529	1	Mon Dec 01 21:35:14 +0000 2014
3	623495536	1	Mon Dec 01 23:55:55 +0000 2014
4	623495537	1	Tue Dec 02 00:06:05 +0000 2014

	text
0	WTF MY BATTERY WAS 31% ONE SECOND AGO AND NOW ...
1	@apple Contact sync between Yosemite and iOS8 ...
2	WARNING IF YOU BUY AN IPHONE 5S UNLOCKED FROM ...
3	@Apple, For the love of GAWD, CENTER the '1'on...
4	i get the storage almost full notification lit...

Bag of Word

```
vectorizer = CountVectorizer()
X_train_counts = vectorizer.fit_transform(X_train)
X_test_counts = vectorizer.transform(X_test)
df_train = pd.DataFrame(X_train_counts.toarray(),
    columns=vectorizer.get_feature_names_out())
df_test = pd.DataFrame(X_test_counts.toarray(),
    columns=vectorizer.get_feature_names_out())

print("Bag of Words for Training Data:")
print(df_train.values[:5])
print("\nBag of Words for Testing Data:")
print(df_test.values[:5])
```

Bag of Words for Training Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words for Testing Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

TF_IDF Bag Of word

```
tfidf_vectorizer = TfidfVectorizer()
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
X_test_tfidf = tfidf_vectorizer.transform(X_test)
df_train_tfidf = pd.DataFrame(X_train_tfidf.toarray(),
    columns=tfidf_vectorizer.get_feature_names_out())
df_test_tfidf = pd.DataFrame(X_test_tfidf.toarray(),
    columns=tfidf_vectorizer.get_feature_names_out())
print("Bag of Words based on TF-IDF for Training Data:")
print(df_train_tfidf.values[:5])
print("\nBag of Words based on TF-IDF for Testing Data:")
print(df_test_tfidf.values[:5])
```

Bag of Words based on TF-IDF for Training Data:

```
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
```

```
[0. 0. 0. ... 0. 0. 0.]
```

Bag of Words based on TF-IDF for Testing Data:

```
[[0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]
```

N_Grams

Unigrams

```
unigram_vectorizer = CountVectorizer(ngram_range=(1, 1))  
X_train_unigrams = unigram_vectorizer.fit_transform(X_train)  
X_test_unigrams = unigram_vectorizer.transform(X_test)  
df_train_unigrams = pd.DataFrame(X_train_unigrams.toarray(),  
                                  columns=unigram_vectorizer.get_feature_names_out())  
df_test_unigrams = pd.DataFrame(X_test_unigrams.toarray(),  
                                 columns=unigram_vectorizer.get_feature_names_out())  
print("Bag of Words based on unigrams for Training Data:")  
print(df_train_unigrams.values[:5])  
print("\nBag of Words based on unigrams for Testing Data:")  
print(df_test_unigrams.values[:5])
```

Bigrams

```
bigram_vectorizer = CountVectorizer(ngram_range=(2, 2))  
X_train_bigrams = bigram_vectorizer.fit_transform(X_train)  
X_test_bigrams = bigram_vectorizer.transform(X_test)  
df_train_bigrams = pd.DataFrame(X_train_bigrams.toarray(),  
                                 columns=bigram_vectorizer.get_feature_names_out())  
df_test_bigrams = pd.DataFrame(X_test_bigrams.toarray(),  
                                columns=bigram_vectorizer.get_feature_names_out())  
print("\nBag of Words based on bigrams for Training Data:")  
print(df_train_bigrams.values[:5])  
print("\nBag of Words based on bigrams for Testing Data:")  
print(df_test_bigrams.values[:5])
```

Trigrams

```
trigram_vectorizer = CountVectorizer(ngram_range=(3, 3))  
X_train_trigrams = trigram_vectorizer.fit_transform(X_train)  
X_test_trigrams = trigram_vectorizer.transform(X_test)  
df_train_trigrams = pd.DataFrame(X_train_trigrams.toarray(),  
                                  columns=trigram_vectorizer.get_feature_names_out())  
df_test_trigrams = pd.DataFrame(X_test_trigrams.toarray(),  
                                 columns=trigram_vectorizer.get_feature_names_out())  
print("\nBag of Words based on trigrams for Training Data:")  
print(df_train_trigrams.values[:5])  
print("\nBag of Words based on trigrams for Testing Data:")  
print(df_test_trigrams.values[:5])
```

Bag of Words based on unigrams for Training Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words based on unigrams for Testing Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words based on bigrams for Training Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words based on bigrams for Testing Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words based on trigrams for Training Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

Bag of Words based on trigrams for Testing Data:

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

```
results = {'Classifier': [], 'Vectorizer': [], 'Accuracy': [],
           'Classification Report': []}
```

```
vectorizers = {
    'Bag of Words (Raw Counts)': CountVectorizer(),
    'Bag of Words (TF-IDF)': TfidfVectorizer(),
    'Unigrams': CountVectorizer(ngram_range=(1, 1)),
    'Bigrams': CountVectorizer(ngram_range=(2, 2)),
```

```

    'Trigrams': CountVectorizer(ngram_range=(3, 3)),
}

```

Nave bise

```

classifiers = {
    'Naïve Bayes': MultinomialNB(),
}
for clf_name, clf in classifiers.items():
    for vec_name, vectorizer in vectorizers.items():
        X_train_vec = vectorizer.fit_transform(X_train)
        X_test_vec = vectorizer.transform(X_test)
        clf.fit(X_train_vec, y_train)
        y_pred = clf.predict(X_test_vec)
        accuracy = accuracy_score(y_test, y_pred)
        classification_rep = classification_report(y_test, y_pred)
        results['Classifier'].append(clf_name)
        results['Vectorizer'].append(vec_name)
        results['Accuracy'].append(accuracy)
        results['Classification Report'].append(classification_rep)

classifier_name = 'Naïve Bayes'
for i in range(len(results['Classifier'])):
    print("\nResults for", results['Classifier'][i], "with",
results['Vectorizer'][i], ":")
    print(f"Accuracy: {results['Accuracy'][i]}")
    print(f"Classification Report:\n{results['Classification Report']
[i]}")

    # Use the correct key for the classifier
    X_train_vec = vectorizers[results['Vectorizer']
[i]].fit_transform(X_train)
    X_test_vec = vectorizers[results['Vectorizer']
[i]].transform(X_test)
    classifiers[classifier_name].fit(X_train_vec, y_train)
    y_pred = classifiers[classifier_name].predict(X_test_vec)

    cm = confusion_matrix(y_test, y_pred, labels=class_names)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.title(f'Confusion Matrix for {classifier_name} with
{results["Vectorizer"][i]}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()

```


Results for Naïve Bayes with Bag of Words (Raw Counts) :

Accuracy: 0.7262210796915167

Classification Report:

	precision	recall	f1-score	support
1	0.66	0.80	0.72	264
3	0.78	0.81	0.79	428
5	0.67	0.11	0.19	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.73	778
macro avg	0.53	0.43	0.43	778
weighted avg	0.71	0.73	0.70	778

Confusion Matrix for Naïve Bayes with Bag of Words (Raw Counts)




Results for Naïve Bayes with Bag of Words (TF-IDF) :

Accuracy: 0.7519280205655527

Classification Report:

	precision	recall	f1-score	support
1	0.80	0.70	0.75	264
3	0.73	0.92	0.81	428
5	1.00	0.07	0.13	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.75	778
macro avg	0.63	0.42	0.42	778
weighted avg	0.77	0.75	0.71	778




Results for Naïve Bayes with Unigrams :

Accuracy: 0.7262210796915167

Classification Report:

	precision	recall	f1-score	support
1	0.66	0.80	0.72	264
3	0.78	0.81	0.79	428
5	0.67	0.11	0.19	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.73	778
macro avg	0.53	0.43	0.43	778
weighted avg	0.71	0.73	0.70	778




Results for Naïve Bayes with Bigrams :

Accuracy: 0.7287917737789203

Classification Report:


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	1	0.73	0.69	0.71	264
	3	0.74	0.88	0.80	428
	5	0.40	0.14	0.20	73
not_relevant		0.00	0.00	0.00	13
accuracy				0.73	778
macro avg		0.47	0.43	0.43	778
weighted avg		0.70	0.73	0.70	778



Results for Naïve Bayes with Trigrams :					
Accuracy: 0.6619537275064268					
Classification Report:					
	precision	recall	f1-score	support	
1	0.72	0.44	0.55	264	
3	0.65	0.91	0.76	428	

5	0.50	0.10	0.16	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.66	778
macro avg	0.47	0.36	0.37	778
weighted avg	0.65	0.66	0.62	778



Logistic Regression

```

classifiers = {
    'Logistic Regression': LogisticRegression(),
}
results = {'Classifier': [], 'Vectorizer': [], 'Accuracy': [],
           'Classification Report': []}
for clf_name, clf in classifiers.items():

```

```

for vec_name, vectorizer in vectorizers.items():
    X_train_vec = vectorizer.fit_transform(X_train)
    X_test_vec = vectorizer.transform(X_test)
    clf.fit(X_train_vec, y_train)
    y_pred = clf.predict(X_test_vec)
    accuracy = accuracy_score(y_test, y_pred)
    classification_rep = classification_report(y_test, y_pred)
    results['Classifier'].append(clf_name)
    results['Vectorizer'].append(vec_name)
    results['Accuracy'].append(accuracy)
    results['Classification Report'].append(classification_rep)

for i in range(len(results['Classifier'])):
    print("\nResults for", results['Classifier'][i], "with",
    results['Vectorizer'][i], ":")
    print(f"Accuracy: {results['Accuracy'][i]}")
    print(f"Classification Report:\n{results['Classification Report']
[i]}")
    X_train_vec = vectorizers[results['Vectorizer']
[i]].fit_transform(X_train)
    X_test_vec = vectorizers[results['Vectorizer']
[i]].transform(X_test)
    classifiers[results['Classifier'][i]].fit(X_train_vec, y_train)
    y_pred = classifiers[results['Classifier'][i]].predict(X_test_vec)
    cm = confusion_matrix(y_test, y_pred, labels=class_names)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.title(f'Confusion Matrix for {results["Classifier"][i]} with
{results["Vectorizer"][i]}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()

```


Results for Logistic Regression with Bag of Words (Raw Counts) :

Accuracy: 0.7493573264781491

Classification Report:

	precision	recall	f1-score	support
1	0.77	0.72	0.75	264
3	0.75	0.88	0.81	428
5	0.54	0.21	0.30	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.75	778
macro avg	0.51	0.45	0.46	778
weighted avg	0.72	0.75	0.73	778

Confusion Matrix for Logistic Regression with Bag of Words (Raw Counts)




Results for Logistic Regression with Bag of Words (TF-IDF) :

Accuracy: 0.7403598971722365

Classification Report:

	precision	recall	f1-score	support
1	0.76	0.70	0.73	264
3	0.73	0.89	0.80	428
5	0.69	0.15	0.25	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.74	778
macro avg	0.55	0.43	0.44	778
weighted avg	0.73	0.74	0.71	778

Confusion Matrix for Logistic Regression with Bag of Words (TF-IDF)




Results for Logistic Regression with Unigrams :

Accuracy: 0.7493573264781491

Classification Report:

	precision	recall	f1-score	support
1	0.77	0.72	0.75	264
3	0.75	0.88	0.81	428
5	0.54	0.21	0.30	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.75	778
macro avg	0.51	0.45	0.46	778
weighted avg	0.72	0.75	0.73	778




Results for Logistic Regression with Bigrams :

Accuracy: 0.7043701799485861

Classification Report:

	precision	recall	f1-score	support
1	0.67	0.65	0.66	264
3	0.72	0.86	0.78	428
5	0.75	0.12	0.21	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.70	778
macro avg	0.54	0.41	0.41	778
weighted avg	0.69	0.70	0.67	778




Results for Logistic Regression with Trigrams :

Accuracy: 0.6516709511568124

Classification Report:

	precision	recall	f1-score	support
1	0.78	0.35	0.48	264
3	0.63	0.95	0.75	428
5	0.80	0.11	0.19	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.65	778
macro avg	0.55	0.35	0.36	778
weighted avg	0.68	0.65	0.60	778



Randome Forest

```

classifiers = {
    'Random Forest': RandomForestClassifier(),
}
results = {'Classifier': [], 'Vectorizer': [], 'Accuracy': [],
           'Classification Report': []}
for clf_name, clf in classifiers.items():
    for vec_name, vectorizer in vectorizers.items():
        X_train_vec = vectorizer.fit_transform(X_train)
        X_test_vec = vectorizer.transform(X_test)
        clf.fit(X_train_vec, y_train)
        y_pred = clf.predict(X_test_vec)
        accuracy = accuracy_score(y_test, y_pred)
        classification_rep = classification_report(y_test, y_pred)
        results['Classifier'].append(clf_name)
        results['Vectorizer'].append(vec_name)

```



```

        results['Accuracy'].append(accuracy)
        results['Classification Report'].append(classification_rep)
for i in range(len(results['Classifier'])):
    clf_name = results['Classifier'][i]
    vec_name = results['Vectorizer'][i]
    X_train_vec = vectorizers[vec_name].fit_transform(X_train)
    X_test_vec = vectorizers[vec_name].transform(X_test)
    classifiers[clf_name].fit(X_train_vec, y_train)
    y_pred = classifiers[clf_name].predict(X_test_vec)
    cm = confusion_matrix(y_test, y_pred, labels=class_names)
    print("\nResults for", clf_name, "with", vec_name, ":")
    print(f"Accuracy: {results['Accuracy'][i]}")
    print(f"Classification Report:\n{results['Classification Report']
[i]}")
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.title(f'Confusion Matrix for {clf_name} with {vec_name}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()

```


Results for Random Forest with Bag of Words (Raw Counts) :

Accuracy: 0.7416452442159382

Classification Report:

	precision	recall	f1-score	support
1	0.75	0.71	0.73	264
3	0.74	0.89	0.81	428
5	0.77	0.14	0.23	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.74	778
macro avg	0.56	0.43	0.44	778
weighted avg	0.73	0.74	0.71	778

Confusion Matrix for Random Forest with Bag of Words (Raw Counts)




Results for Random Forest with Bag of Words (TF-IDF) :

Accuracy: 0.7480719794344473

Classification Report:

	precision	recall	f1-score	support
1	0.74	0.72	0.73	264
3	0.75	0.89	0.81	428
5	0.69	0.15	0.25	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.75	778
macro avg	0.55	0.44	0.45	778
weighted avg	0.73	0.75	0.72	778




Results for Random Forest with Unigrams :

Accuracy: 0.7416452442159382

Classification Report:

	precision	recall	f1-score	support
1	0.75	0.70	0.73	264
3	0.74	0.88	0.81	428
5	0.68	0.18	0.28	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.74	778
macro avg	0.54	0.44	0.45	778
weighted avg	0.73	0.74	0.72	778

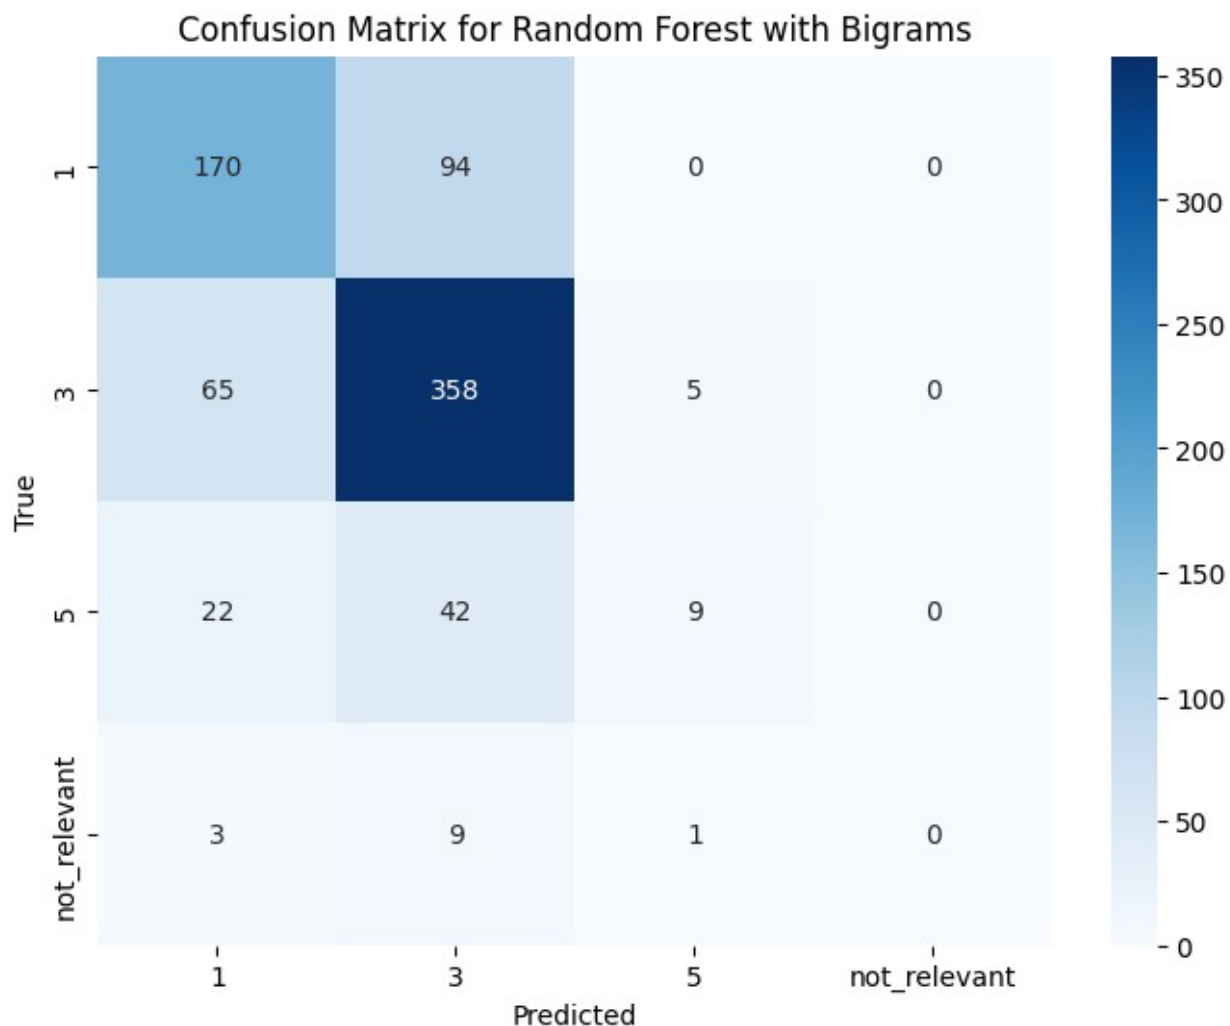


Results for Random Forest with Bigrams :

Accuracy: 0.6979434447300771

Classification Report:

	precision	recall	f1-score	support
1	0.65	0.68	0.67	264
3	0.72	0.83	0.77	428
5	0.78	0.10	0.17	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.70	778
macro avg	0.54	0.40	0.40	778
weighted avg	0.69	0.70	0.67	778

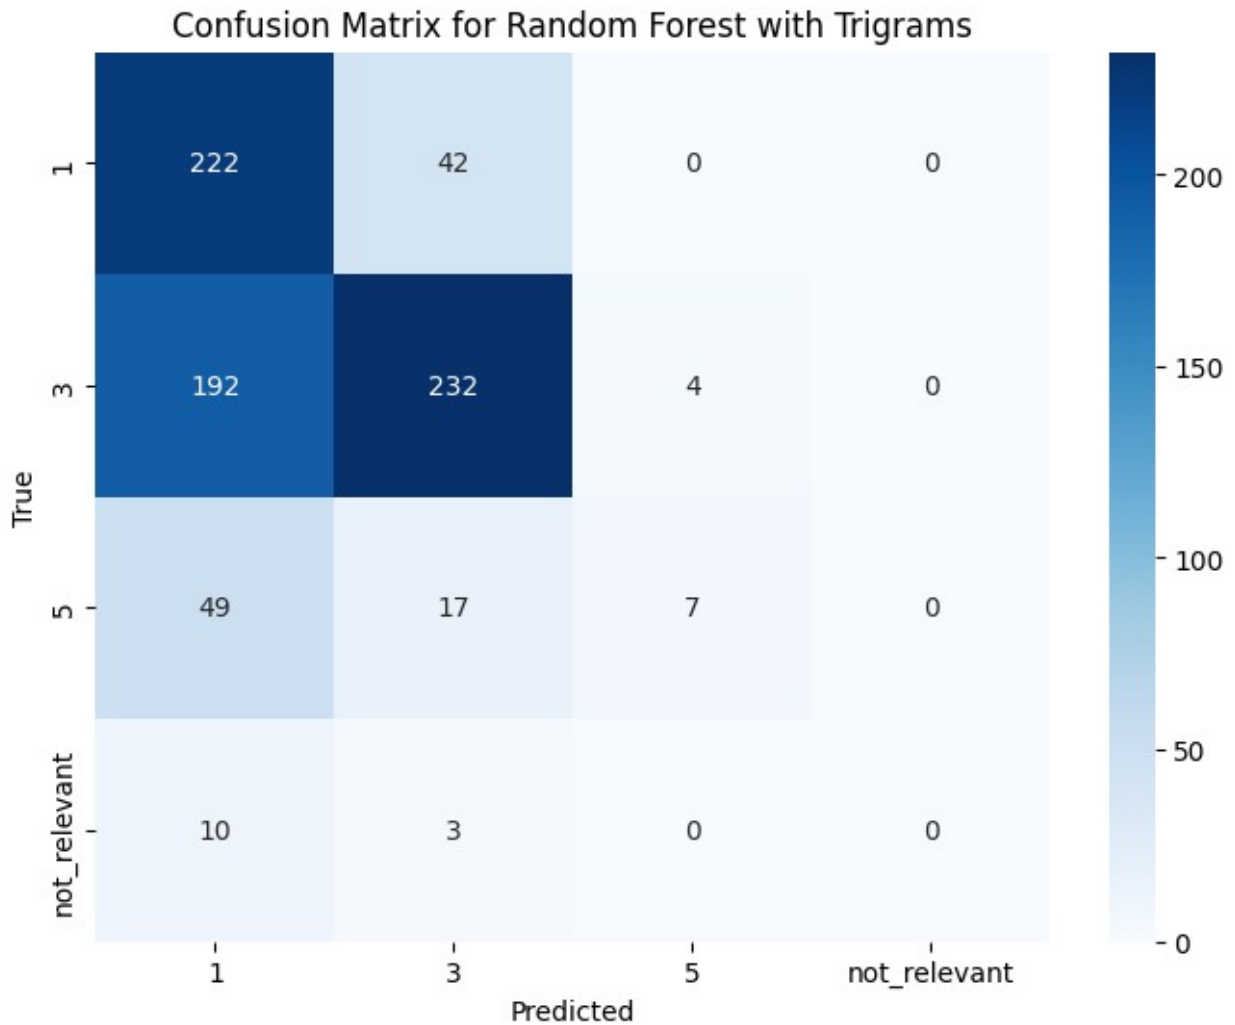


Results for Random Forest with Trigrams :

Accuracy: 0.5912596401028277

Classification Report:

	precision	recall	f1-score	support
1	0.47	0.84	0.60	264
3	0.79	0.54	0.64	428
5	0.64	0.10	0.17	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.59	778
macro avg	0.47	0.37	0.35	778
weighted avg	0.65	0.59	0.57	778



SVM

```
classifiers = {
    'SVM': SVC(),
}
results = {'Classifier': [], 'Vectorizer': [], 'Accuracy': [],
           'Classification Report': []}
for clf_name, clf in classifiers.items():
    for vec_name, vectorizer in vectorizers.items():
        X_train_vec = vectorizer.fit_transform(X_train)
        X_test_vec = vectorizer.transform(X_test)
        clf.fit(X_train_vec, y_train)
        y_pred = clf.predict(X_test_vec)
        accuracy = accuracy_score(y_test, y_pred)
        classification_rep = classification_report(y_test, y_pred)
        results['Classifier'].append(clf_name)
        results['Vectorizer'].append(vec_name)
```

```

        results['Accuracy'].append(accuracy)
        results['Classification Report'].append(classification_rep)

for i in range(len(results['Classifier'])):
    print("\nResults for", results['Classifier'][i], "with",
results['Vectorizer'][i], ":")
    print(f"Accuracy: {results['Accuracy'][i]}")
    print(f"Classification Report:\n{results['Classification Report']
[i]}")
    clf_name = results['Classifier'][i]
    vec_name = results['Vectorizer'][i]
    X_train_vec = vectorizers[vec_name].fit_transform(X_train)
    X_test_vec = vectorizers[vec_name].transform(X_test)
    classifiers[clf_name].fit(X_train_vec, y_train)
    y_pred = classifiers[clf_name].predict(X_test_vec)
    cm = confusion_matrix(y_test, y_pred, labels=class_names)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.title(f'Confusion Matrix for {clf_name} with {vec_name}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()


```

Results for SVM with Bag of Words (Raw Counts) :

Accuracy: 0.7352185089974294

Classification Report:

	precision	recall	f1-score	support
1	0.74	0.68	0.71	264
3	0.73	0.90	0.81	428
5	0.88	0.10	0.17	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.74	778
macro avg	0.59	0.42	0.42	778
weighted avg	0.74	0.74	0.70	778




Results for SVM with Bag of Words (TF-IDF) :

Accuracy: 0.7506426735218509

Classification Report:

	precision	recall	f1-score	support
1	0.76	0.71	0.73	264
3	0.74	0.90	0.81	428
5	0.85	0.15	0.26	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.75	778
macro avg	0.59	0.44	0.45	778
weighted avg	0.75	0.75	0.72	778




Results for SVM with Unigrams :

Accuracy: 0.7352185089974294

Classification Report:

	precision	recall	f1-score	support
1	0.74	0.68	0.71	264
3	0.73	0.90	0.81	428
5	0.88	0.10	0.17	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.74	778
macro avg	0.59	0.42	0.42	778
weighted avg	0.74	0.74	0.70	778




Results for SVM with Bigrams :

Accuracy: 0.6979434447300771

Classification Report:

	precision	recall	f1-score	support
1	0.64	0.69	0.67	264
3	0.73	0.83	0.78	428
5	0.86	0.08	0.15	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.70	778
macro avg	0.56	0.40	0.40	778
weighted avg	0.70	0.70	0.67	778




Results for SVM with Trigrams :

Accuracy: 0.6015424164524421

Classification Report:

	precision	recall	f1-score	support
1	0.89	0.15	0.25	264
3	0.58	0.99	0.73	428
5	0.86	0.08	0.15	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.60	778
macro avg	0.58	0.30	0.28	778
weighted avg	0.70	0.60	0.50	778



Perceptron

```
classifiers = {
    'Perceptron': Perceptron(),
}
results = {'Classifier': [], 'Vectorizer': [], 'Accuracy': [],
           'Classification Report': []}
for clf_name, clf in classifiers.items():
    for vec_name, vectorizer in vectorizers.items():
        X_train_vec = vectorizer.fit_transform(X_train)
        X_test_vec = vectorizer.transform(X_test)
        clf.fit(X_train_vec, y_train)
        y_pred = clf.predict(X_test_vec)
        accuracy = accuracy_score(y_test, y_pred)
        classification_rep = classification_report(y_test, y_pred)
        results['Classifier'].append(clf_name)
        results['Vectorizer'].append(vec_name)
```

```

        results['Accuracy'].append(accuracy)
        results['Classification Report'].append(classification_rep)
for i in range(len(results['Classifier'])):
    print("\nResults for", results['Classifier'][i], "with",
results['Vectorizer'][i], ":")
    print(f"Accuracy: {results['Accuracy'][i]}")
    print(f"Classification Report:\n{results['Classification Report']
[i]}")
    X_train_vec = vectorizers[results['Vectorizer']
[i]].fit_transform(X_train)
    X_test_vec = vectorizers[results['Vectorizer']
[i]].transform(X_test)
    classifiers[results['Classifier'][i]].fit(X_train_vec, y_train)
    y_pred = classifiers[results['Classifier'][i]].predict(X_test_vec)
    cm = confusion_matrix(y_test, y_pred, labels=class_names)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.title(f'Confusion Matrix for {results["Classifier"][i]} with
{results["Vectorizer"][i]}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()


```

Results for Perceptron with Bag of Words (Raw Counts) :

Accuracy: 0.7339331619537275

Classification Report:

	precision	recall	f1-score	support
1	0.80	0.66	0.72	264
3	0.73	0.89	0.80	428
5	0.49	0.23	0.31	73
not_relevant	0.17	0.08	0.11	13
accuracy			0.73	778
macro avg	0.55	0.46	0.49	778
weighted avg	0.72	0.73	0.72	778




Results for Perceptron with Bag of Words (TF-IDF) :

Accuracy: 0.7005141388174807

Classification Report:

	precision	recall	f1-score	support
1	0.66	0.80	0.72	264
3	0.77	0.73	0.75	428
5	0.51	0.29	0.37	73
not_relevant	0.08	0.08	0.08	13
accuracy			0.70	778
macro avg	0.51	0.47	0.48	778
weighted avg	0.70	0.70	0.69	778




Results for Perceptron with Unigrams :

Accuracy: 0.7339331619537275

Classification Report:

	precision	recall	f1-score	support
1	0.80	0.66	0.72	264
3	0.73	0.89	0.80	428
5	0.49	0.23	0.31	73
not_relevant	0.17	0.08	0.11	13
accuracy			0.73	778
macro avg	0.55	0.46	0.49	778
weighted avg	0.72	0.73	0.72	778




Results for Perceptron with Bigrams :

Accuracy: 0.6516709511568124

Classification Report:

	precision	recall	f1-score	support
1	0.63	0.67	0.65	264
3	0.71	0.73	0.72	428
5	0.42	0.23	0.30	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.65	778
macro avg	0.44	0.41	0.42	778
weighted avg	0.64	0.65	0.64	778



Results for Perceptron with Trigrams :

Accuracy: 0.6208226221079691

Classification Report:

	precision	recall	f1-score	support
1	0.62	0.46	0.53	264
3	0.64	0.81	0.72	428
5	0.42	0.19	0.26	73
not_relevant	0.00	0.00	0.00	13
accuracy			0.62	778
macro avg	0.42	0.37	0.38	778
weighted avg	0.60	0.62	0.60	778

Confusion Matrix for Perceptron with Trigrams

